
Dog Breed Identification: An application of Convolutional Neural Network and Transfer Learning on Inception and EfficientNet

Jaki Tang
Master of Science in Analytics
University of Chicago
jakitang@uchicago.edu

Jolina Shen
Master of Science in Analytics
University of Chicago
jshen7@uchicago.edu

William Dai
Master of Science in Analytic
University of Chicago
wd9@uchicago.edu

Abstract

Dogs have lots of breeds. Some of the breeds are very similar to each other and most of the people cannot distinguish. In this project, we want to investigate how current image classification neural networks perform on classifying dog breeds. We plan to start from applying either baseline CNN or pre-trained models like Efficient Net B0 and Inception V3. Then, we optimize the models through adding layers like drop out or batch normalization and through hyper parameter tuning of the optimizers. Without image preprocessing, we achieved train accuracy 96.91% and validation accuracy 42.42% for Inception, and train accuracy 99.94% and validation accuracy 64.25% for Efficient Net. After image preprocessing, the validation accuracy increases to 78.96% for Inception and 79.49% for Efficient Net.

20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46

Table of Contents

Abstract	1
Exploratory Data Analysis	3
Data Count and Description	3
Visualization of the Raw Data	3
Univariate Analysis	4
Model Building	4
Deep Learning Algorithm Description	5
Inception V3	5
Efficient Net B0	6
Model Performance Evaluation and Metrics Selection	7
Loss Function	7
Accuracy Metric	8
Model Management	8
Model Architecture Deployment	8
Inception	8
Efficient Net	8
Activation Function	8
Optimizer	8
Model Maintenance and Parameter Update	8
Model Results	9
Conclusion	10
Discussion of Findings	10
Acknowledgments	10
References and Resources	10

1 Exploratory Data Analysis

Our dataset is a subset of ImageNet [1] containing 10222 images for training and 10000 images for testing.

1.1 Data Count and Description

The dataset contains 10222 images of 120 dog breeds in train dataset with a csv file containing each image's label information. Average images per dog breed is 85 images. The dataset also contains 10000 unlabeled images as test dataset. There are no missing values in the dataset.

1.2 Visualization of the Raw Data

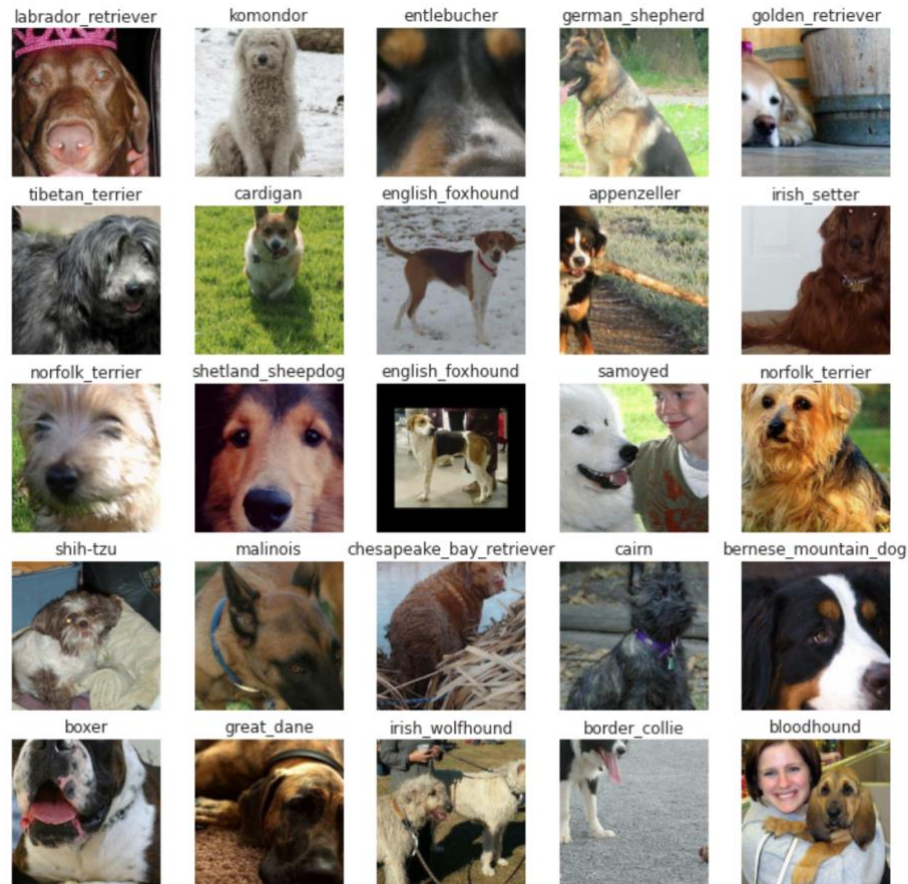


Figure 1: Raw Data Visualization

Figure 1 is a visualization showing some of the training dataset's images. Each image of dog is labeled with its corresponding breed.

1.3 Univariate Analysis




Figure 2: Train and Test Data Distribution

Above is the visualization of a bar chart showing the distribution of images per dog breed in the training dataset. From the univariate analysis, the maximum images in one dog breed are 126 for scottish_deerhound and the minimum images in one dog breed is 66 for eskimo_dog or briard. The standard deviation is 13.24. Therefore, the train dataset is in a good distribution without selection bias.

2 Model Building

We first implemented the baseline CNN and make a test run on our dataset to decide the direction of our next step.

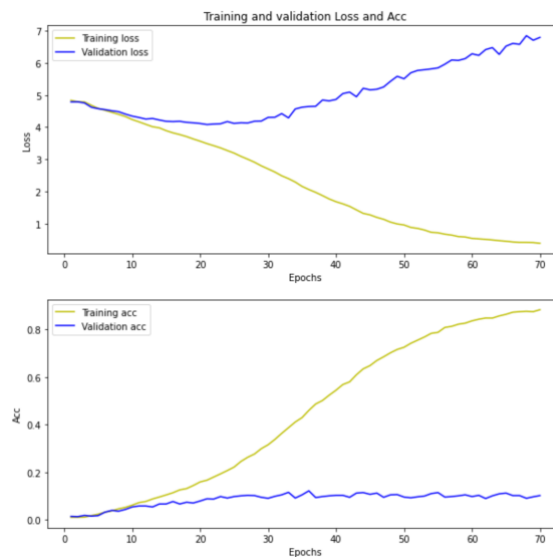


Figure 3: Train and Validation Loss for Baseline CNN

The Figure 3 shows that baseline CNN cannot handle the images' complexity in our dataset.

2.1 Deep Learning Algorithm Description

Against the baseline CNN. We noticed that the baseline CNN cannot capture all the features in our dataset. For example, the lack of width and depth coefficient in baseline CNN fails to extract necessary features from the images. Therefore, we choose to build one Efficient Net B0 and one Inception V3. Both Inception and Efficient Net are deeper and wider than Baseline CNN thus should give us better results.

2.1.1 Inception V3

We choose InceptionV3. Inception marks an important milestone in the development of CNNs. Inception adopts Inception module which contains multiple kernels with different sizes so that objects of various sizes can be captured from the previous layer. This feature is important to our dataset as dog in each image has different size so a fixed size kernel cannot handle all our data properly. We choose V3 over V1 because V3 Inception factorized those 3x3 and 5x5 kernels so that the computational efficiency is increased [2].




Figure 4: How Inception Module Works for Multiple Kernel Sizes

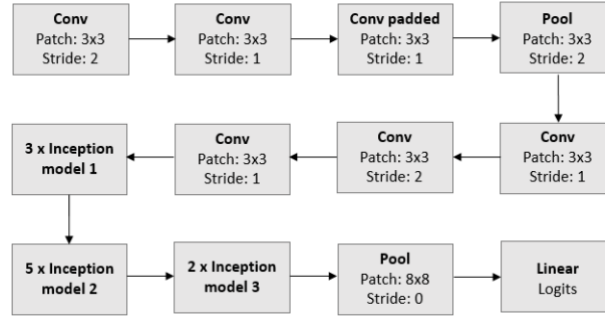


Figure 5: Inception V3 Design

In addition, Inception V3 adds more regularization to the loss function which could prevent overfitting. Because many dog breeds are very similar, we don't want our model to be too confident in classification.




Figure 6: Train and Validation Loss for Inception before Preprocessing




Figure 7: Train and Validation Loss for Inception after Preprocessing

The Figure 5 and 6 shows that our inception model encounter overfitting in the beginning and it cannot be solved simply through model hyper parameter tuning. Therefore, we add normalization to rescale images' pixels by 255 in the preprocessing step. Then, the overfitting issue is much less severe as the validation accuracy increases from 42.42% to 78.96%.

2.1.2 Efficient Net B0

We choose EfficientNet because it is pretrained on more than a million images from the ImageNet database. EfficientNet is newly developed in 2019. It refines the scaling of all dimensions of the neural network like width, depth, and image resolution against the current CNNs to improve the overall performance [6]. The technique used by EfficientNet is compound coefficient. The compound scaling combines the design of width scaling, depth scaling and higher resolution scaling in one design with fixed coefficients. The three coefficients control widths, depths and resolution, and will be balanced with a constant ratio.

The normal grid search achieves total FLOPs (Floating Point Operations Per Second) around 2 but EfficientNet increases this by 2 of the computational resources' square [5].

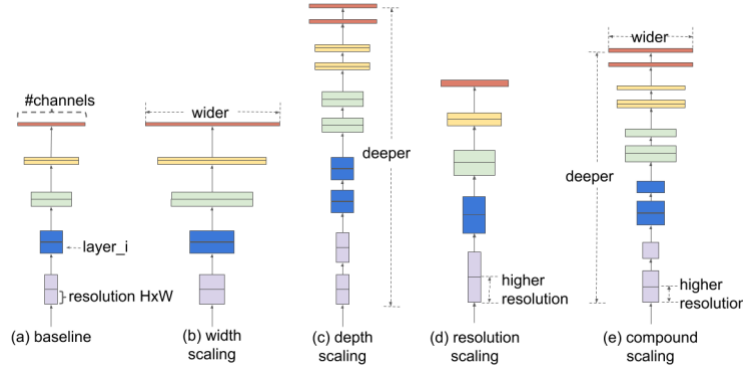


Figure 8: Scaling Method of Efficient Net

The design of EfficientNet is developed by performing a neural architecture search using the AutoML MNAS framework. It adds the FLOPs increase layers to the original design and it balances accuracy and time complexity by penalizing heavy computations because of the large scaling process [6].




Figure 9: Design of Efficient Net

130 We choose B0 as our version because the accuracy of B0 is closer to InceptionV3 and we want to
131 perform some model tuning on our own




Figure 10: Train and Validation Loss for Efficient Net before Preprocessing

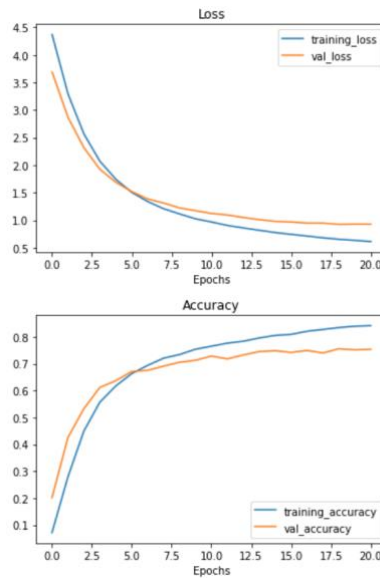


Figure 11: Train and Validation Loss for Efficient Net after Preprocessing

The Figure 10 and 11 shows that our Efficient Net model encounter overfitting in the beginning. However, the overfitting is not as worse as Inception. We think it may be the reason that the Efficient Net is pretrained on ImageNet. Therefore, the overfitting is not as worse as Inception. In addition, we add normalization to rescale images' pixels by 255 in the preprocessing step. Then, the overfitting issue is much less severe as the validation accuracy increases from 64.25% to 79.49%.

2.2 Model Performance Evaluation and Metrics Selection

The model performance evaluation is done through categorical_crossentropy and accuracy.

2.2.1 Loss Function

We use categorical_crossentropy over binary_crossentropy because our encoded output label

has 120 classes. And the output label is assigned one-hot category encoding value in form of 0s and 1s [3]. The binary crossentropy would fail to capture the losses between our 120 labeled classes.

151

152 **2.2.2 Accuracy Metric**

153 We use accuracy as our metrics. If we choose to use binary accuracy, the 120 classes of our
154 label would result in high binary accuracy for all our predictions. Accuracy can better reflect
155 our performance.

156 Above Figures (Figure 5, 6, 10, and 11) show how our models perform. The results shown
157 some overfitting even when we add dropout layers, batch normalization, and tune the
158 optimizer with learning rate and decay.

159 Therefore, we choose to preprocess our image with rescale equal to 1.0/255.0 so that all the
160 images are normalized. The results of models processing normalized images are shown
161 below. The fixes some of the overfitting.

162

163 **3 Model Management**

164

165 **3.1 Model Architecture Deployment**

166

167 **3.1.1 Inception**

168 Based on the Inception, we also add one `global_average_pooling` layer, two dropout layers,
169 two batch normalization layers, and two dense layers.

170

171 **3.1.2 Efficient Net**

172 Based on the Efficient Net, we add one `global_average_pooling` layer as well as one dense
173 layer. The `global_average_pooling` layer is followed by a dropout layer and a batch
174 normalization layer to reduce the overfitting problem.

175

176 **3.1.3 Activation Function**

177 Both models use softmax as activation function instead of sigmoid because softmax is better
178 for multiclass classification as when softmax is used, it increases the probability of one class
179 and decreases the total probability of all other classes. This feature fits our data.

180

181 **3.1.4 Optimizer**

182 We also choose Adam as our optimizer over SGD because Adam implicitly performs
183 coordinate-wise gradient clipping and can hence, unlike SGD, tackle heavy-tailed noise [4].
184

185 **3.2 Model Maintenance and Parameter Update**

186 Through our model tuning, we focus on the below parts: adding dense layers with dropout
187 and batch normalization, tuning dropout rate, adding weight regularization on dense layers,
188 and optimizing optimizers through learning rate and decay.

189 The dropout rate is tested from 0.25 to 1 in step size of 0.25 for Inception and 0.1 to 0.5 in
190 step size of 0.1 for Efficient Net. The smaller dropout rate is, the better it can help decrease
191 the influence of overfitting.

192 The weight regularization is added as l1 and l2 regularization to both kernel and bias. L2
193 regularization works better than l1 regularization.

194 Learning rate is tested as 0.0001, 0.0005 and 0.001. Smaller learning rate helps with
195 overfitting, but larger learning rate helps improve computation complexity. Therefore, we
196 switch back to 0.001 when we add image preprocessing.

197 Decay is tested as 1e-6 and 1e-5. Both have improvements on overfitting, but difference is

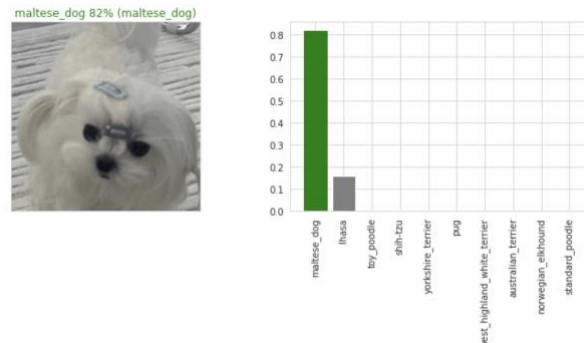
198 subtle.

199

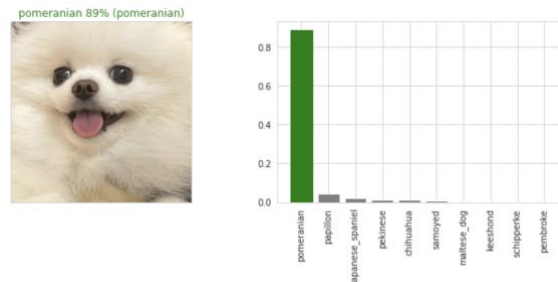
200 **3.3 Model Results**

201 Four sample prediction from the efficient Net is appended below. As we can see that our
202 model still has a rather strict requirements in the image's quality because it is hard to
203 distinguish features like texture or relative size of the dog.


204 However, even the final prediction is not accurate, but we still have the right breed or similar
205 breed across the confidence intervals.
206



207

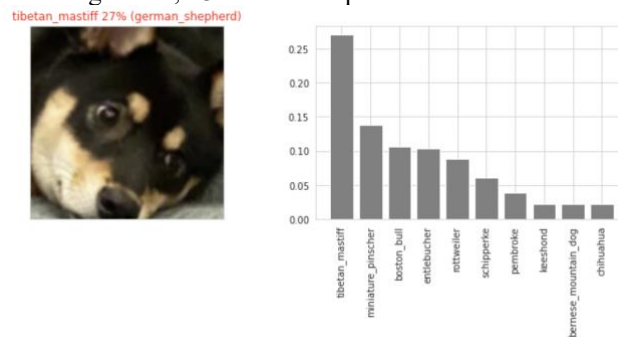


208



209

210 Figure 12, 13 & 14: Sample Good Prediction



211

212 Figure 15: Sample Bad Prediction

213

214 **4 Conclusion**

215

216 **4.1 Discussion of Findings**

217 In this project, we want to take use of CNN's image recognition ability to identify dog breeds
218 for us. We first fit and test our dataset with a baseline CNN but the result on validation
219 dataset is not improving against each epoch. Also, the validation loss keeps going higher.
220 Therefore, we believe the complexity of baseline CNN is not high enough to handle this
221 dataset. We then try to train this dataset with Inception V3 and EfficientNet B0. Both these
222 two neural networks have a deeper and wider design to handle more features from dog
223 images. The initial result without any image normalization shows signals of overfitting as
224 Inception has train accuracy 96.91% and validation accuracy 42.42%, and EfficientNet has
225 train accuracy 99.94% and validation accuracy 64.25%. After image preprocessing with
226 normalization, the validation accuracy increases to 78.96% for Inception and 79.49% for
227 Efficient Net. Therefore, we find that dog breeds identification is a much more complex
228 problem against classification between different type of animals like dog vs cat. Both
229 Inception and EfficientNet achieve a good result but we may investigate further in the future
230 for more improvements.

231

232 **4.2 Acknowledgments**

233 We would like to thank Professor Ashis Pujari for his teaching and answering our problems.

234

235 **4.3 References and Resources**

236 [1] Olga Russakovsky*, Jia Deng*, Hao Su, Jonathan Krause, Sanjeev Satheesh, Sean Ma, Zhiheng
237 Huang, Andrej Karpathy, Aditya Khosla, Michael Bernstein, Alexander C. Berg and Li Fei-Fei. (* =
238 equal contribution) ImageNet Large Scale Visual Recognition Challenge. IJCV, 2015.

239 [2] Raj, B., 2022. A Simple Guide to the Versions of the Inception Network. [online] Available at:
240 <[https://towardsdatascience.com/a-simple-guide-to-the-versions-of-the-inception-network-7fc52b8632](https://towardsdatascience.com/a-simple-guide-to-the-versions-of-the-inception-network-7fc52b863202)
241 02> [Accessed 30 May 2022].

242 [3] Kumar, A., 2022. [online] Available at: <[https://vitalflux.com/keras-categorical-cross-entropy-loss-](https://vitalflux.com/keras-categorical-cross-entropy-loss-function/)
243 [function/](https://vitalflux.com/keras-categorical-cross-entropy-loss-function/)> [Accessed 30 May 2022].

244 [4] Bushaev, Vitaly. "Adam-Latest Trends in Deep Learning Optimization." Medium, Towards Data
245 Science, 24 Oct. 2018, [https://towardsdatascience.com/adam-latest-trends-in-deep-learning-](https://towardsdatascience.com/adam-latest-trends-in-deep-learning-optimization-6be9a291375c)
246 [optimization-6be9a291375c](https://towardsdatascience.com/adam-latest-trends-in-deep-learning-optimization-6be9a291375c).

247 [5] EfficientNet paper — Tan, M., & Le, Q. v. (2019). EfficientNet: Rethinking model scaling for
248 convolutional neural networks. 36th International Conference on Machine Learning, ICML 2019,
249 2019-June.

250 [6] "EfficientNet: Improving Accuracy and Efficiency through AutoML and Model Scaling." Google
251 AI Blog, 29 May 2019, <https://ai.googleblog.com/2019/05/efficientnet-improving-accuracy-and.html>.

252 [7] "Dog Breed Identification." Kaggle.com, 29 Sept. 2017, [www.kaggle.com/competitions/dog-breed-](http://www.kaggle.com/competitions/dog-breed-identification/overview)
253 [identification/overview](http://www.kaggle.com/competitions/dog-breed-identification/overview). Accessed 30 May 2022.

254 [8] SAXENA, ATRI. "Using Tensorflow 2.x to Classify Breeds of Dogs." Kaggle.com, 30 Dec. 2021,
255 www.kaggle.com/code/atrisaxena/using-tensorflow-2-x-to-classify-breeds-of-dogs. Accessed 30 May
256 2022.

257 [9] VARGHESE, JITHIN. "Accuracy: 89.71% InceptionV3 Keras." Kaggle.com, 10 May 2021,
258 www.kaggle.com/code/jithinanievarghese/accuracy-89-71-inceptionv3-keras. Accessed 30 May 2022.

259 [10] "Dog Breed CNN." Kaggle.com, www.kaggle.com/code/anu332/dog-breed-cnn.